

强化学习基本原理及编程实现06：提升学习效率的方法

郭宪

2019.11.3

人工智能学院

College of Artificial Intelligence




南开大学
Nankai University



提升学习效率的方法

1. N 步时间差分方法
2. 资格迹方法-- $TD(\lambda)$
3. off-policy






第一部分：N 步时间差分方法（第七章）

值函数估计过程

Backup 值:

表格型值函数估计



$$\text{DP } Q_{\pi}(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a \sum_{a' \in A} \pi(a'|s') Q_{\pi}(s', a')$$




值函数估计过程

Backup 值:

$$\text{DP } Q_{\pi}(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a \sum_{a' \in A} \pi(a'|s') Q_{\pi}(s', a')$$



函数逼近: $\hat{q}(s, a, \theta)$



训练目标: $\arg \min_{\theta} \in (q(s, a) - \hat{q}(s, a, \theta))^2$

强化学习: 在线学习

值函数的评估

增量式MC方法估计值函数：

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

$\Rightarrow G_t$ 是值函数 $v_\pi(S_t)$ 的无偏估计

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$

最简单的时间差分学习算法：TD(0)

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

\Rightarrow 真实的TD目标 $R_{t+1} + \gamma v_\pi(S_{t+1})$ 是无偏估计，
但 $R_{t+1} + \gamma V(S_{t+1})$ 是有偏估计

$R_{t+1} + \gamma V(S_{t+1})$ 称为TD目标

TD目标 $R_{t+1} + \gamma v_\pi(S_{t+1})$ 的方差比MC的返回值 G_t 要小很多。因为MC的返回值依赖于很多随机动作，转移概率和回报。TD目标仅依赖于一个随机动作，转移概率和回报。

N步预测

值函数估计:

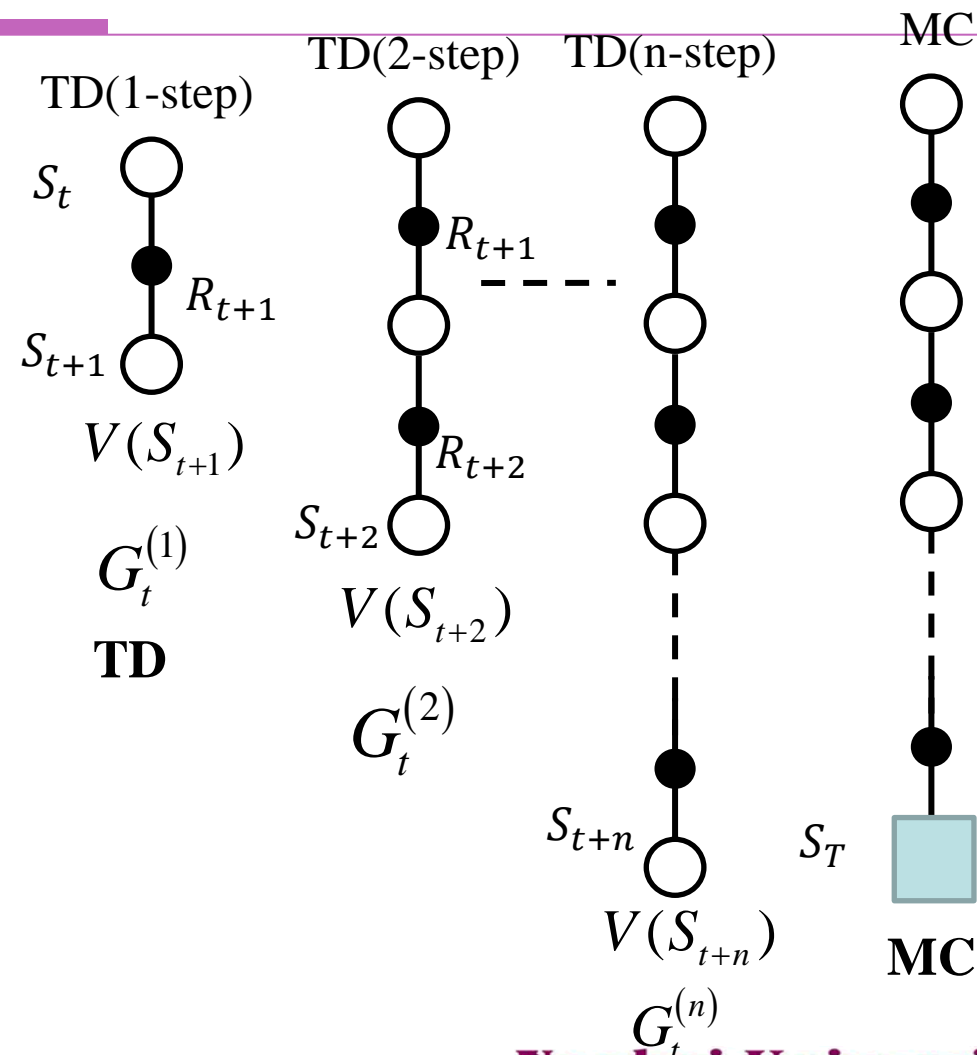
$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

$$G_t^{(1)} = R_{t+1} + \gamma V(S_{t+1})$$

$$G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V(S_{t+1})$$

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$

$$\text{MC}: G_t^{(mc)} = R_{t+1} + \gamma R_{t+2} + \dots$$



N步时间差分算法

n -step TD for estimating $V \approx v_\pi$

Input: a policy π

Algorithm parameters: step size $\alpha \in (0, 1]$, a positive integer n

Initialize $V(s)$ arbitrarily, for all $s \in \mathcal{S}$

All store and access operations (for S_t and R_t) can take their index mod $n + 1$

Loop for each episode:

Initialize and store $S_0 \neq \text{terminal}$

$T \leftarrow \infty$

Loop for $t = 0, 1, 2, \dots$:

 If $t < T$, then:

 Take an action according to $\pi(\cdot | S_t)$

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being updated)


 If $\tau \geq 0$:

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$


 If $\tau + n < T$, then: $G \leftarrow G + \gamma^n V(S_{\tau+n})$ ($G_{\tau:\tau+n}$)

$V(S_\tau) \leftarrow V(S_\tau) + \alpha [G - V(S_\tau)]$

Until $\tau = T - 1$




单步预测




两步预测

N步时间差分算法




N步Sarsa


Path taken




Action values increased
by one-step Sarsa



Action values increased
by 10-step Sarsa



N步Sarsa比单步Sarsa要快，多步Sarsa能学到更多的知识



第二部分：资格迹方法-- $TD(\lambda)$

N步预测

值函数估计:

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$


$$G_t^{(1)} = R_{t+1} + \gamma V(S_{t+1})$$

$$G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V(S_{t+2})$$


$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$

$$G_t^\lambda = (1-\lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

$$= (1-\lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} G_t^{(n)} + \lambda^{T-t-1} G_t$$



$TD(\lambda)$ 前向视角




$$V(s_t) \leftarrow V(s_t) + \alpha (G_t^{(\lambda)} - V(s_t))$$

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

为了得到当前值函数的估计需要将来的回报和值函数，因此像蒙特卡罗方法一样，只有整个实验结束后，才能计算得到。

$TD(\lambda)$ 后向视角



当前的TD偏差: $\delta_t = R_{t+1} + \gamma V(s_{t+1}) - V(s_t)$

对于每个状态 s ，值函数的改变量为: $V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$ ← 适合度轨迹

适合度轨迹定义:
$$E_t(s) = \begin{cases} \gamma \lambda E_{t-1}(s), & \text{if } s \neq s_t \\ \gamma \lambda E_{t-1}(s) + 1, & \text{if } s = s_t \end{cases}$$



资格迹的理解：表格型

访问状态流：

s_0													
-------	--	--	--	--	--	--	--	--	--	--	--	--	--

资格迹：

1	0	0	0	0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---

访问状态流：

s_0	s_1												
-------	-------	--	--	--	--	--	--	--	--	--	--	--	--

资格迹：

$\lambda\gamma$	1	0	0	0	0	0	0	0	0	0	0	0	0
-----------------	---	---	---	---	---	---	---	---	---	---	---	---	---



$TD(\lambda), TD(0), TD(1)$

当 $\lambda=0$, 只有当前状态值更新:

$$v(S_t) \leftarrow V(S_t) + \alpha \delta_t$$

当 $\lambda=1$, 状态 s 值函数的总更新与MC等价

$$\begin{aligned} & \delta_t + \gamma \delta_{t+1} + \gamma^2 \delta_{t+2} + \dots + \gamma^{T-1-t} \delta_{T-1} \\ &= R_{t+1} + \gamma V(S_{t+1}) - V(S_t) \\ &+ \gamma R_{t+2} + \gamma^2 V(S_{t+2}) - \gamma V(S_{t+1}) \\ &+ \gamma^2 R_{t+3} + \gamma^3 V(S_{t+3}) - \gamma^2 V(S_{t+2}) \\ &\vdots \\ &+ \gamma^{T-1-t} R_T + \gamma^{T-t} V(S_T) - \gamma^{T-1-t} V(S_{T-1}) \\ &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-1-t} R_T - V(S_t) \\ &= G_t - V(S_t) \end{aligned}$$

对于一般的 λ 前向视角偏差等价于后向视角偏差

$$\begin{aligned} G_t^\lambda - V(S_t) &= -V(S_t) + (1-\lambda) \lambda^0 (R_{t+1} + \gamma V(S_{t+1})) \\ &+ (1-\lambda) \lambda^1 (R_{t+1} + \gamma R_{t+2} + \gamma^2 V(S_{t+2})) \\ &+ (1-\lambda) \lambda^2 (R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 V(S_{t+3})) + \dots \\ &= -V(S_t) + (\gamma \lambda)^0 (R_{t+1} + \gamma V(S_{t+1}) - \gamma \lambda V(S_{t+1})) \\ &+ (\gamma \lambda)^1 (R_{t+2} + \gamma V(S_{t+2}) - \gamma \lambda V(S_{t+2})) \\ &+ (\gamma \lambda)^2 (R_{t+3} + \gamma V(S_{t+3}) - \gamma \lambda V(S_{t+3})) \\ &+ \dots \\ &= (\gamma \lambda)^0 (R_{t+1} + \gamma V(S_{t+1}) - V(S_t)) \\ &+ (\gamma \lambda)^1 (R_{t+2} + \gamma V(S_{t+2}) - V(S_{t+1})) \\ &+ (\gamma \lambda)^2 (R_{t+3} + \gamma V(S_{t+3}) - V(S_{t+2})) + \dots \\ &= \delta_t + \gamma \lambda \delta_{t+1} + (\gamma \lambda)^2 \delta_{t+2} + \dots \end{aligned}$$

Sarsa(λ)算法

1. 初始化 $Q(s, a), \forall s \in S, a \in A(s)$, 给定参数 α, γ

2. Repeat:

行动策略和评估策略都是 ϵ 贪婪策略

给定起始状态 s , 并根据 ϵ 贪婪策略在状态 s 选择动作 a , 对所有的 $s \in S, a \in A(s), E(s, a) = 0$

Repeat (对于一幕的每一步)

(a) 根据 ϵ 贪婪策略在状态 s 选择动作 a , 得到回报 r 和下一个状态 s' , 在状态 s' 根据 ϵ 贪婪策略得到动作 a'

(b) $\delta \leftarrow r + \gamma Q(s', a') - Q(s, a), E(s, a) \leftarrow E(s, a) + \delta$

(c) 对所有的 $s \in S, a \in A(s): Q(s, a) \leftarrow Q(s, a) + \alpha \delta E(s, a), E(s, a) \leftarrow \gamma \lambda E(s, a)$

(d) $s = s', a = a'$

Until s 是终止状态

Until 所有的 $Q(s, a)$ 收敛

3. 输出最终策略: $\pi(s) = \arg \max_a Q(s, a)$



资格迹的理解：函数逼近型

资格迹为与权重维数相同的向量，为短期记忆，持续的时间少于一幕的长度，其作用是辅助整个学习过程，具体过程为：

$$z_{-1} \doteq \mathbf{0}$$


$$z_t \doteq \gamma \lambda z_{t-1} + \nabla \hat{v}(S_t, w_t)$$

时间差分误差为：

$$\delta_t \doteq R_{t+1} + \gamma \hat{v}(S_{t+1}, w_t) - \hat{v}(S_t, w_t)$$


权重更新方法为：

$$w_{t+1} \doteq w_t + \alpha \delta_t z_t$$



第三部分： off-policy

重要性采样



重要性采样

重要性采样求积分：

$$\begin{aligned} E[f] &= \int f(z) p(z) dz \\ &= \int f(z) \frac{p(z)}{q(z)} q(z) dz \\ &\approx \frac{1}{N} \sum_n \frac{p(z^n)}{q(z^n)} f(z^n), z^n \sim q(z) \end{aligned}$$

定义重要性权重： $\omega^n = p(z^n) / q(z^n)$

普通的重要性采样求积分： $E[f] = \frac{1}{N} \sum_n \omega^n f(z^n)$

重要性采样积分：无偏估计，但方差无穷大

减小方差的方法：加权重要性采样求积分

$$E[f] \approx \sum_{n=1}^N \frac{\omega^n}{\sum_{m=1}^N \omega^m} f(z^n)$$

MC 重要性采样

在策略 π 下, t 时刻后轨迹的概率为:

$$\Pr(A_t, S_{t+1}, \dots, S_T) = \prod_{k=t}^{T-1} \pi(A_k | S_k) p(S_{k+1} | S_k, A_k)$$

在目标策略和行为策略下, 每个回报都使用概率进行加权

$$\rho_t^T = \frac{\prod_{k=t}^{T-1} \pi(A_k | S_k) p(S_{k+1} | S_k, A_k)}{\prod_{k=t}^{T-1} \mu(A_k | S_k) p(S_{k+1} | S_k, A_k)} = \prod_{k=t}^{T-1} \frac{\pi(A_k | S_k)}{\mu(A_k | S_k)}$$

普通重要性采样, 值估计:

时间 t 后的第一次终止时刻 从 t 到 $T(t)$ 的返回值

$$V(s) = \frac{\sum_{t \in \mathcal{T}(s)} \rho_t^{T(t)} G_t}{|\mathcal{T}(s)|}$$

状态 s 被访问过的所有时刻的集合

s s

$t = 1 \ 2 \ 3 \ 4 \ 5 \ 6 \ 7 \ 8 \ 9 \ 10 \ 11 \ 12 \ 13 \ 14 \ 15 \ 16 \ 17 \ 18 \ 19$

$$\mathcal{T}(s) = \{4, 15\} \quad T(4) = 7, T(15) = 19$$

加权重要性采样, 值估计:

$$V(s) = \frac{\sum_{t \in \mathcal{T}(s)} \rho_t^{T(t)} G_t}{\sum_{t \in \mathcal{T}(s)} \rho_t^{T(t)}}$$

Off-policy every visit MC

初始化, 对于所有的

$$s \in S, a \in A(s):$$

$$Q(s, a) \leftarrow \text{任意}$$

$$C(s, a) \leftarrow 0$$

$$\pi(s) \leftarrow \text{相对于 } Q \text{ 的贪婪策略}$$

Repeat forever:

利用软策略 μ 产生一次实验:

$$S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T, S_T$$

$$G \leftarrow 0$$

$$W \leftarrow 1$$

For $t = T-1, T-2, \dots$ down to 0:

$$G \leftarrow \gamma G + R_{t+1}$$

$$C(S_t, A_t) \leftarrow C(S_t, A_t) + W$$

策略评估


$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{W}{C(S_t, A_t)} [G - Q(S_t, A_t)]$$

$$\pi(S_t) \leftarrow \arg \max_a Q(S_t, a)$$

策略改善

如果 $A_t \neq \pi(S_t)$ 则退出for循环

$$W \leftarrow W \frac{1}{\mu(A_t | S_t)}$$





Off-policy n-step Sarsa (7.3节)

Off-policy n -step Sarsa for estimating $Q \approx q_*$ or q_π

Input: an arbitrary behavior policy b such that $b(a|s) > 0$, for all $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize $Q(s, a)$ arbitrarily, for all $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize π to be greedy with respect to Q , or as a fixed given policy

Algorithm parameters: step size $\alpha \in (0, 1]$, a positive integer n

All store and access operations (for S_t , A_t , and R_t) can take their index mod $n + 1$

Loop for each episode:

Initialize and store $S_0 \neq \text{terminal}$

Select and store an action $A_0 \sim b(\cdot|S_0)$

$T \leftarrow \infty$

Loop for $t = 0, 1, 2, \dots$:

 If $t < T$, then:

 Take action A_t

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then:

$T \leftarrow t + 1$

 else:

 Select and store an action $A_{t+1} \sim b(\cdot|S_{t+1})$

$\tau \leftarrow t - n + 1$ (τ is the time whose estimate is being updated)

 If $\tau \geq 0$:

$\rho \leftarrow \prod_{i=\tau+1}^{\min(\tau+n-1, T-1)} \frac{\pi(A_i|S_i)}{b(A_i|S_i)} \quad (\rho_{\tau+1:t+n-1})$

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

 If $\tau + n < T$, then: $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n}) \quad (G_{\tau:\tau+n})$

$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha \rho [G - Q(S_\tau, A_\tau)]$

 If π is being learned, then ensure that $\pi(\cdot|S_\tau)$ is greedy wrt Q

Until $\tau = T - 1$

